

XIYUAN YANG

yangxiyuan@sjtu.edu.cn ◇ Shanghai Jiao Tong University ◇ Shanghai, China

GitHub: <https://github.com/xiyuanyang-code>

Online Resume: <https://xiyuanyang-code.github.io/cv/>

EDUCATION

Shanghai Jiao Tong University

School of Artificial Intelligence

GPA: 4.1/4.3 (Ranked 2 out of 66)

Score: 93.9/100

Shanghai, China

09/2024 - present

Six Scholarships:

- National Scholarship (2025, 2026) (2 of 66)
- SJTU Undergraduate Excellence Scholarship (2024)
- Zhiyuan Honor Scholarship (2024, 2025)
- "Han Ying Ju Hua" Scholarship (2025, 15 per year.)

High graded courses:

- Comprehensive Programming: 100
- Prob. and Statistics (Honor): 100
- Algorithm Design: 100
- Numerical Analysis: 100
- Linear Algebra (Honor): 98
- C++ Programming Design (Honor): 98

26 courses achieved A/A+ (16 A+), including all major-related courses.

RESEARCH

Current Research Directions:

- Scalable Agentic Data Synthesis & Training.** Generating high-quality instances and agent trajectories for mid- and post-training, with an emphasis on long-horizon tasks.
- Recursive Self-Improvement (RSI).** Developing training- and inference-time methods for efficient, iterative model improvement and automated post-training.

Research Experience:

- MAGIC LAB** (2024.12 - present): Undergraduate researcher advised by Prof. [Siheng Chen](#).
Search-Agent, Agentic Benchmark, Auto-PostTraining (RSI).
- IQuest Research** (2026.7 - present): Research intern (IQuest-Coder LLMs, Post-Training).
Long-Horizon Agentic Instance Synthesis & Mid/Post-Training Trajectory Data Recipe.

PUBLICATIONS & TECH PROJECTS

Tech-Report IQuest-Q1: Advancing Foundation Capabilities for Agentic CLI Systems

[IQuestLab/IQuest-Q1](#) [IQuest-Q1 Model Page](#) [Project Page](#)

Contributed to **IQuest-Q1**, a 320B-parameter Mixture-of-Experts model with 15B active parameters per token, designed for agentic coding, reasoning, and multi-step tool use in CLI environments. The model is trained on synthetic medium- and long-horizon tasks using multi-harness reinforcement learning and multi-teacher on-policy distillation (MOPD).

RSI-Master: Structuring Experiments to Guide Autonomous Model Improvement

Incoming submission for ICLR 2027 (Co-first author)

[arXiv: 2609.35561](#) [DorothyDUUU/RSI-Master](#)

Pioneered **RSI-Master**, which combines an Experiment OS with Reviewer-Guided Research Orchestration to prevent reward hacking and strategy lock-in in autonomous model development. On PostTrainBench with Qwen3-4B-Base, it achieved 54.49 versus 46.53 for the strongest agent baseline with a 0.0% hacking rate; at 35B, it surpassed the human-developed Instruct model on LiveCodeBench-v6 (41.21 vs. 37.36) and SciCode.

DataMaster: Data-Centric Autonomous AI Research
Incoming submission for ICLR 2027 (Co-first author)

✂ arXiv: 2605.10906 📄 sjtu-sai-agents/DataMaster

Proposed and developed DataMaster automating external data discovery, selection, and cleaning to enhance ML and post training tasks. Designed a system combining tree-structured search, shared data pool to navigate open-ended search spaces and carry global memory across branches. Successfully boosted the MLE-Bench Lite medal rate by 32.27% and surpassed the instruct model on PostTrainBench (GPQA: 31.02% vs. 30.35%).

ICML-2026 MCP-Persona: Benchmarking LLM Agents on Real-World Personal Applications via Environment Simulation

Co-first author ✂ arXiv: 2606.02470 📄 wwh0411/MCP-Persona

Pioneered **MCP-Persona**, the first benchmark for evaluating LLM agents on **personalized** Model Context Protocol (MCP) tools. MCP-Persona captures real-world challenges of tools interacting with individual accounts and local databases across diverse platform including Reddit, Xiaohongshu, Lark, and Slack. Extensive evaluations of state-of-the-art agents reveal substantial deficiencies in personalized tool utilization.

ICLR-2026 InfoMosaic-Bench: Evaluating Multi-Source Information Seeking in Tool-Augmented Agents
3rd Author ✂ arXiv: 2510.02271 📄 DorothyDUUU/Info-Mosaic

Pioneered **InfoMosaic-Bench**, the first benchmark dedicated to multi-source information seeking in LLM agents, demonstrating the critical need to integrate general-purpose web search with thousands of domain-specific tools via **MCP** (Model Context Protocol). Experiments revealed that web information alone is insufficient (GPT-5 Acc: 38.2%) for information seeking due to incorrect tool usage/selection.

Tech-Report AppCopilot: Toward General, Accurate, Long-Horizon, and Efficient Mobile Agent
📄 OpenBMB/AppCopilot (300+ stars.)

Developed **AppCopilot**, a multi-modal, multi-agent on-device assistant that addresses core challenges in the mobile-agent landscape. AppCopilot significantly improves generalization, on-screen interaction accuracy, and long-horizon task completion, while optimizing performance and efficiency on mobile devices.

OPEN SOURCE & SKILLS

Technical Writing. I maintain a **technical blog** (<https://xiyuanyang-code.github.io>) with 150+ articles and 480k+ words on computer science and AI. I also maintain **Frontier LLM Notes** (<https://xiyuanyang-code.github.io/llm/>), covering model architectures, training recipes, LLM infrastructure and source-code analysis, and agentic systems.

Skills. LLM (Model Architectures, Training, Agentic Harness), Programming Languages (Python, C++), Training Infra (slime, megatron, harbor). English CET-6: 584.

AWARDS AND CERTIFICATIONS

Challenge Cup Competition for "Unveiling the List and Appointing the Best" Defender	2025
Honorable Mention in MCM	2025, 2026
National Scholarship (2 of 66)	2025, 2026
Zhiyuan Honor Scholarship	2024, 2025
"Han Ying Ju Hua" Scholarship (15 per year)	2025
SJTU Undergraduate Excellence Scholarship	2024